

Fluentd meets Unicode Windows EventLog

Fluentd meetup 2019

ClearCode Inc.

Hiroshi Hatake

Agenda

- Motivation
- About winevt_c
- Unicode Character handling
- Using ANSI code page issues
- Unicode Testing
- Benchmark
- Throughput Benchmark
- Conclusion

Agenda

- **Motivation**
- About winevt_c
- Unicode Character handling
- Using ANSI code page issues
- Unicode Testing
- Benchmark
- Throughput Benchmark
- Conclusion

Motivation

- in_windows_eventlog has some issues...
 - 🤔 Unicode character handling. Sometimes garbage chracters are generated.
 - 🤔 Memory consumption in flood of windows event
 - 🤔 Sometimes it causes SEGV
 - 😊 CPU spike when resuming operation
 - 🙌 At least one event should exist in the listening channel on starting to listen. Otherwise, nothing to be read
- And they are caused by dependent gem which is named win32-eventlog

Next Topic

- Motivation
- **About winevt_c**
- Unicode Character handling
- Using ANSI code page issues
- Unicode Testing
- Benchmark
- Throughput Benchmark
- Conclusion

winevt_c (new gem): Code examples

Just querying for specified channel

```
require 'winevt'

@query = Winevt::EventLog::Query.new(
  "Application", "[System[(Level <= 3) and TimeCreated[timediff(@SystemTime) <= 86400000]]]"
)

@query.each do |eventlog, message, string_inserts|
  puts ({eventlog: eventlog, data: message})
end
```

winevt_c (new gem): Code examples

Update bookmark for querying channel

```
require 'winevt'

@query = Winevt::EventLog::Query.new(
  "Application", "[System[Level <= 3) and TimeCreated[timediff(@SystemTime) <= 86400000]]")
@bookmark = Winevt::EventLog::Bookmark.new
@query.each do |xml|
  @bookmark.update(@query)
end

puts @bookmark.render
```

winevt_c (new gem): Code examples

Subscribe channel

```
require 'winevt'



@subscribe = Winevt::EventLog::Subscribe.new
@subscribe.tail = true
@subscribe.subscribe(
  "Security", "[System[(Level <= 4) and TimeCreated[timediff(@SystemTime) <= 86400000]]]"
)
while true do
  @subscribe.each do |eventlog, message, string_inserts|
    puts ({eventlog: eventlog, data: message})
  end
  sleep(1)
end
```


winevt_c (new gem)

- It solves win32-eventlog issues
 - 😊 Improve Unicode character handling.
 - 😊 It doesn't cause SEGV on the same situation
 - 😊 CPU spike when resuming operation is declined
 - 😊 Reduce memory consumption in flood of windows event
 - This issue still exists but it is reduced memory consumption
 - 😊 At least one event should exist in the listening channel on starting to listen.
 - Empty channel can also subscribe. The older one will be staled.

winevt_c (new gem)

The relationship of plugins and gems in this talk



Next Topic

- Motivation
- About winevt_c
- **Unicode Character handling**
- Using ANSI code page issues
- Unicode Testing
- Benchmark
- Throughput Benchmark
- Conclusion

Unicode Character handling

- What is *Unicode*?

In Windows context, it means *UTF-16*.

In Ruby C extension context, it means *UTF-8*.

Unicode Character handling:

What is the difference between ANSI and Unicode?

- In Windows, *ANSI* means current code page
 - In Japanese Edition Windows, it is CP932 (Windows-31J).
 - **-A** suffixed API uses ANSI character encoding
- In Windows, *Unicode* means UTF-16
 - **-W** suffixed API uses UTF-16 character encoding
 - **PWSTR** and such **W** contained typed API arguments also use UTF-16 character encoding

Unicode Character handling

- We need to convert from UTF-16 to target character encoding
 - In this case, target encoding is *UTF-8*
- But, win32-eventlog gem uses `OpenEventLogA`, `ReadEventLogA` (ANSI version)
 - To handle Unicode characters correctly, we need to use `OpenEventLogW`, `ReadEventLogW` (Unicode version)
 - win32-eventlog gem development is inactive in recent days.
 - Unicode version patch exists, but it have not been merged in....

Next Topic

- Motivation
- About winevt_c
- Unicode Character handling
- **Using ANSI code page issues**
- Unicode Testing
- Benchmark
- Throughput Benchmark
- Conclusion

Using ANSI code page issues

- On Japanese Edition Windows' default code page can handle...
 - Alphabets
 - Greek letters
 - Cyrillic alphabets
 - Hiragana, Katakana
 - JIS level 1 and 2 Kanji sets (Chinese Characters)
- But other characters cannot handle with cp932 (In Japanese Edition Windows)

Using ANSI code page issues: UTF-8 contains more characters!

- UTF-8 can also handles...
 - Alphabets
 - Greek letters
 - Cyrillic alphabets
 - Hiragana, Katakana
 - JIS level 1 and 2 Kanji set (Chinese Characters)
- And...
 - diacritical mark (such as umlaut in German: ä, ö, ü)
 - Hebrew, Arabic, Devanagari (Hindi)
 - South East Asia Characters (Thai, Laotian... etc.)
 - And **Emoji!!** 😎

Using ANSI code page issues: Solution



- We decide to develop the brand new gem which is named winevt_c.
 - 🥰 It uses new Windows API that is defined in <[winevt.h](#)>
 - 😊 The new API provides bookmark which is used to resume operation
 - 😊 Unicode API
- But this gem is written in C and C++
 - 🛠️ Users need to build C/C++ extension code by themselves
 - 😊 Current RubyInstaller bundles MSYS2 system. Users can use gcc and g++ after MSYS2 installation which is kicked by RubyInstaller.

Next Topic

- Motivation
- About winevt_c
- Unicode Character handling
- Using ANSI code page issues
- **Unicode Testing**
- Benchmark
- Throughput Benchmark
- Conclusion

Unicode Testing: Environment

- Windows 10 Home 1903 64bit
 - Japanese Edition (cp932, Windows-31J)
- Writing Windows EventLog
 - [Benchmark tool](#) written in C#
- Terminal:
 - PowerShell Core 6 on Windows Terminal(Preview)
 - Used Windows Terminal Profile is [here](#)



Unicode Testing: Writing Events in .NET ([picked up](#))

```
for (int i = 0; i < totalEvents / 10; i++)
{
    if (i % 10 == 0)
    {
        Console.WriteLine(String.Format("{0, 8}", i * 10));
        Task.Run(() => MonitorProcesses(counter));
    }

    // Write an informational entry to the event log.
    benchLog.WriteEntry(String.Format("Writing to event log. {0} times.", i)); // Alphabets
    benchLog.WriteEntry("ⓕⓁⓊⓎⓈⓉⓂ"); // Non-ASCII symbols
    benchLog.WriteEntry("日本語による説明"); // Japanese
    benchLog.WriteEntry("สวัสดีจาก Fluentd!"); // Thai
    benchLog.WriteEntry("Привет, от Fluentd."); // Cyrillic
    benchLog.WriteEntry("Γεια σου, από την Fluentd."); // Greek letters
    benchLog.WriteEntry("مرحبًا ، من Fluentd."); // Arabic alphabets
    benchLog.WriteEntry("हाय, Fluentd से!"); // Devanagari
    benchLog.WriteEntry("We ♥ Fluentd!(●'◡'●)"); // Unicod-ish Kaomoji
    benchLog.WriteEntry("Logging is fun! 😊😄😁😂😃"); // Emoji
    Thread.Sleep(waitMSec);
}
```

Unicode Testing: Writing Events

```
PS> EventLogbencher.exe -w 10 -t 10
```

- 10 Events Written into Benchmark channel

Unicode Testing: Configuration (old plugin)

```
1 <source>
2   @type windows_eventlog
3   @id old-winevtlog
4   tag raw.winevt
5   channels ["Benchmark"]
6   read from head true
7   from_encoding Windows-31j
8   encoding UTF-8
9   # parse_description true
10  <storage>
11    @type local
12    persistent true
13    path ./tmp/storage-old.json
14  </storage>
15 </source>
16
17 <match **>
18   @type stdout
19 </match>
```

from_encoding/encoding parameters are needed to handle character encoding correctly but still unhandled characters exist.

And using default read_interval: 2s.

Unicode Testing: Configuration (new plugin)

```
1 <source>
2   @type windows_eventlog2
3   @id winevtlog
4   tag raw.winevt
5   channels ["Benchmark"]
6   read_from_head true
7   # parse_description true
8   <storage>
9     @type local
10    persistent true
11    path ./tmp/storage.json
12  </storage>
13 </source>
14
15 <match **>
16   @type stdout
17 </match>
```

No need to specify
from_encoding/encoding
parameters. And new plugin
always handles character
encoding as UTF-8.

And using default
read_interval: 2s.

Unicode Testing: Execution Log (old plugin)

```
Admin: fluent-plugin-windows-eventlog [master] ~ PowerShell 6.2.1 64-bit (48232) x
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "Writing to event log. 27
times.\r\n", "string_inserts": ["Writing to event log. 27 times."]}
2019-07-11 16:26:18.890873000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071484", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "日本語による説明\r\n", "st
ring_inserts": ["日本語による説明"]}
2019-07-11 16:26:18.891192000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071484", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "Привет, от Fluentd.\r\n",
"string_inserts": ["Привет, от Fluentd."]}
2019-07-11 16:26:18.891508000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071486", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "Γεια σου, απ? την Fluentd
.\r\n", "string_inserts": ["Γεια σου, απ? την Fluentd."]}
2019-07-11 16:26:18.891931000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071487", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "??? Fluentd!\r\n",
"string_inserts": ["??? Fluentd!"]}
2019-07-11 16:26:18.892354000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071487", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "??? Fluentd.\r\n",
"string_inserts": ["??? Fluentd."]}
2019-07-11 16:26:18.892677000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071488", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "???, Fluentd ??! \r\n",
"string_inserts": ["???, Fluentd ??!" ]}
2019-07-11 16:26:18.893089000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071489", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "We ? Fluentd!(\u2022'?\u2022)\r\n",
"string_inserts": ["We ? Fluentd!(\u2022'?\u2022)"]}
2019-07-11 16:26:18.893445000 +0900 raw.winevt: {"channel": "benchmark", "record_number": "1071490", "time_generated": "2019-
07-11 16:20:49 +0900", "time_written": "2019-07-11 16:20:49 +0900", "event_id": "0", "event_type": "information", "event_catego
ry": "0", "source_name": "FluentBench", "computer_name": "DESKTOP-G457RDR", "user": "", "description": "Logging is fun! ??????????"
\r\n", "string_inserts": ["Logging is fun! ??????????"]}
```

The following characters are *broken*

- Symbol (F)(L)(T)(D)™
- Thai
- Arabic
- Devanagari (Hindi)
- Unicode contained Kaomoji
- Emoji

Unicode Testing: Execution Log (new plugin)

```

Admin: fluent-plugin-windows-eventlog [master] ~ PowerShell (48232) X
EventRecordID:"1071482","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"Writing to event log. 27 times.,"EventData":["Writing to event log. 27
times."]]
2019-07-11 16:21:31.026289000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071483","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"F L u e n T D ™","EventData":["F L u e n T D ™"]}
2019-07-11 16:21:31.027663 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071484","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"日本語による説明","EventData":["日本語による説明"]}
2019-07-11 16:21:31.029050000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071485","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"สวัสดี จาก Fluentd!"","EventData":["สวัสดี จาก Fluentd!"]}
2019-07-11 16:21:31.030886000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071486","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"Привет, от Fluentd.,"EventData":["Привет, от Fluentd.]}
2019-07-11 16:21:31.032213000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071487","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"Γεια σου, από την Fluentd.,"EventData":["Γεια σου, από την Fluentd.]}
2019-07-11 16:21:31.033554000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071488","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"من Fluentd. من","EventData":["من Fluentd. من"]}
2019-07-11 16:21:31.035122000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071489","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"हाय, Fluentd से!"","EventData":["हाय, Fluentd से!]}
2019-07-11 16:21:31.036541000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071490","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"We ♥ Fluentd!(●'◡'●)"","EventData":["We ♥ Fluentd!(●'◡'●)}
2019-07-11 16:21:31.038058000 +0900 raw.winevt: {"ProviderName":"FluentBench","ProviderGUID":"","EventID":"0","Qualifier
s":"0","Level":"4","Task":"","Opco
EventRecordID:"1071491","ActivityID":"","RelatedActivityID":"","ThreadID":"","Channel":"Benchmark","Computer":"DESKTOP-
G457RDR","UserID":"","Version":"","Description":"Logging is fun! 😊😄😁😂😃"}

```

The following characters are
rendered


- Symbol (F)(L)(T)(D)™
- Thai
- Arabic (*but slightly wrong rendered*)
- Devanagari (Hindi)
- Unicode contained Kaomoji
- Emoji

Next Topic


- Motivation
- About winevt_c
- Unicode Character handling
- Using ANSI code page issues
- Unicode Testing
- **Benchmark**
- Throughput Benchmark
- Conclusion

Benchmark

- Collector Node
 - Windows 10 1809 2 vCPU 4GB Standard SSD
 - [Benchmark tool](#) written in C#
- Aggregator Node
 - Ubuntu 18.04 2 vCPU 4GB Standard SSD
- They are also Azure instances



Collector Node



Aggregator Node

Benchmark: Flow Rate of Events

- 1000000 events total
- About 91 events per seconds

```
PS> EventLogbencher.exe -w 100 -t 10000000
```

- 1 million Events Written into Benchmark channel

Benchmark: Configuration (old)

Collector node

```
<source>
  @type windows_eventlog
  @id old-winevtlog
  tag raw.winevt
  channels ["Benchmark"]
  read_from_head true
  # parse_description true
  <storage>
    @type local
    persistent true
    path ./tmp/storage-old.json
  </storage>
</source>

<match **>
  @type forward
  <server>
    host ["#{ENV['AggregatorServer']}"]
    port 24224
  </server>
  flush_interval 2s
</match>
```

Aggregator node

```
<source>
  @type forward
</source>
<match raw.winevt>
  @type null # or stdout
</match>
```

Benchmark: Configuration (new)

Collector node

Aggregator node


```
<source>
  @type windows_eventlog2
  @id winevtlog
  tag raw.winevt
  channels ["Benchmark"]
  read_from_head true
  # parse_description true
  <storage>
    @type local
    persistent true
    path ./tmp/storage.json
  </storage>
</source>

<match **>
  @type forward
  <server>
    host "#{ENV['AggregatorServer']}"
    port 24224
  </server>
  flush_interval 2s
</match>
```

```
<source>
  @type forward
</source>
<match raw.winevt>
  @type null # or stdout
</match>
```


Benchmark (old plugin)

Flow rate: 91.34361 events per seconds



Benchmark (new plugin)

Flow rate: 91.30634 events per seconds



Benchmark Result: in_windows_eventlog

- Pros
 - 😊 Low CPU usage
- Cons
 - 😞 High memory usage
 - 😞 Incomplete Unicode handling

Benchmark Result: in_windows_eventlog2

- Pros

- 😊 Low memory usage
- 😊 Unicode handling
- 😊 Immediately subscribe channel even if it's empty on subscribe

- Cons

- 😞 Slightly higher CPU usage rather than old plugin's

Next Topic

- Motivation
- About winevt_c
- Unicode Character handling
- Using ANSI code page issues
- Unicode Testing
- Benchmark
- **Throughput Benchmark**
- Conclusion

Throughput Benchmark

- Collector Node
 - Windows 10 1809 2 vCPU 4GB Standard SSD
 - [Benchmark tool](#) written in C#
- Aggregator Node
 - Ubuntu 18.04 2 vCPU 4GB Standard SSD
- They are also Azure instances



Throughput Benchmark

- 500000 events total
- Increase flow rate of events step by step
 - PS> EventLogbencher.exe -w 50 -t 5000000
 - 159.4378 events per seconds
 - PS> EventLogbencher.exe -w 30 -t 5000000
 - 293.4133 events per seconds
 - PS> EventLogbencher.exe -w 20 -t 5000000
 - 314.823 events per seconds
 - PS> EventLogbencher.exe -w 15 -t 5000000
 - 321.7238 events per seconds
 - PS> EventLogbencher.exe -w 10 -t 5000000
 - Stuck 🐞
 - 598.8318 events per seconds
 - **chunk bytes limit exceeds for an emitted event stream** warning is generated from Fluentd....

Throughput Benchmark: Configuration

Collector node


Aggregator node

```
<source>
  @type windows_eventlog2
  @id winevtlog
  tag raw.winevt
  channels ["Benchmark"]
  read_from_head true
  # parse_description true
  <storage>
    @type local
    persistent true
    path ./tmp/storage.json
  </storage>
</source>
```


```
<match **>
  @type forward
  <server>
    host "#{ENV['AggregatorServer']}"
    port 24224
  </server>
  flush_interval 2s
</match>
```

```
<source>
  @type forward
</source>
<match raw.winevt>
  @type null # or stdout
</match>
```


Throughput Benchmark: Result 1




Throughput Benchmark: Result 2



Throughput Benchmark: Result 3



Throughput Benchmark: Result 4



Conclusion

- The new plugin which is named `in_windows_eventlog2` does...
 - Improve Unicode handling
 - Reduce memory consumption
 - Solve CPU spike after resuming operation
- The new plugin might be going to solve...
 - Slightly higher CPU usage than old plugin's
- The new plugin can handle about 300 events per second with default read interval.

Epilogue: Current fluent-plugin-windows-eventlog status

- The new plugin which is named `in_windows_eventlog2`
 - Included fluent-plugin-windows-eventlog v0.3.0
 - We want to hear more user voices and use cases
 - Installation is harder than the older one

Let's enjoy Monitoring Windows EventLog! 😁

Any Questions?